

Assignment 3

Walter Verwer & Bas Machielsen

January 22, 2021

Question 1

1.1

```
data <- tribble(
  ~ color, ~ no_treated, ~ no_contr, ~ avg_treated, ~avg_control,
  "purple", 100, 100, 9, 7,
  "blue", 75, 25, 13, 8,
  "green", 25, 75, 10, 9
)

data %>%
  group_by(color) %>%
  summarize(treatment_effect = avg_treated - avg_control) %>%
  kable(caption = "Treatment effect per color")
```

Table 1: Treatment effect per color

color	treatment_effect
blue	5
green	1
purple	2

1.2

The ATE is defined as $\mathbb{E}[\delta] = \mathbb{E}[Y_1^*] - \mathbb{E}[Y_0^*]$ which are the expectations of the potential outcomes. In general, these two variables are not observed. Under the *random assignment* assumption, we assume that $\mathbb{E}[Y_1^*] = \mathbb{E}[Y_1^*|D = 1]$ and $\mathbb{E}[Y_0^*] = \mathbb{E}[Y_0^*|D = 0]$, which can be estimated by their sample equivalents:

```
data %>%
  summarize(
    e_y1_d_is_1 = (9*100 + 13 * 75 + 10 * 25) / sum(no_treated),
    e_y0_d_is_0 = (7 * 100 + 8 * 25 + 9 * 75) / sum(no_contr)) %>%
  summarize(ate = e_y1_d_is_1 - e_y0_d_is_0) %>%
  kable(caption = "ATE")
```

Table 2: ATE

ate
2.75

1.3

The ATT is defined as $\mathbb{E}[\delta|D = 1] = \mathbb{E}[Y^1|D = 1] - \mathbb{E}[Y^0|D = 1]$. The first term is readily observable. The second term is estimated by us as $\hat{\mathbb{E}}[Y^0|D = 1] = \mathbb{E}[Y^0|D = 0]$. Hence:

```
data_ate <- data %>%
  mutate(n = no_treated + no_contr) %>%
  summarize(e_y1_d_is_1 = (9 * 100 + 13 * 75 + 10 * 25) / sum(no_treated),
            e_y0_d_is_0 = (7 * 100 + 8 * 25 + 9 * 75) / sum(no_contr))

data_ate

## # A tibble: 1 x 2
##   e_y1_d_is_1 e_y0_d_is_0
##   <dbl>      <dbl>
## 1      10.6      7.88
```

```
data_ate %>%
  summarize(att = e_y1_d_is_1 - e_y0_d_is_0) %>%
  kable(caption = "ATT")
```

Table 3: ATT

att
2.75

So the $ATE = ATT$ (because of randomization).

Question 2

2.1

Compute the fraction of students in all three groups (control, low-reward and high-reward) that complete all first-year courses before the start of the second academic year. Show within a table that background characteristics are balanced over the treatment groups.

```
bonus_clean <- bonus %>%
  pivot_longer(cols = c(bonus0, bonus500, bonus1500),
               names_to = "kind_treatment",
               values_to = "treatment") %>%
  filter(treatment == 1)
```

```
bonus_clean %>%
  group_by(kind_treatment) %>%
  summarize(fraction_pass = sum(pass)/n()) %>%
  kable(caption = "Fraction passed per treatment", digits = 3)
```

Table 4: Fraction passed per treatment

kind_treatment	fraction_pass
bonus0	0.195
bonus1500	0.241
bonus500	0.202

```
#Drop the outcome variables
bonus_clean %>%
  group_by(kind_treatment) %>%
  select(-c(pass, stp2001, stp2004, dropout)) %>%
  summarize(across(p0:math,
    list(mean = ~ mean(., na.rm = TRUE),
          sd = ~ sd(., na.rm = TRUE)))) %>%
  pivot_longer(-kind_treatment, names_to = "variable") %>%
  separate(variable, into = c("var", "statistic"), sep = "_") %>%
  pivot_wider(names_from = c(kind_treatment, statistic)) %>%
  kable(caption = "Means and SDs according to treatment", digits = 3)
```

Table 5: Means and SDs according to treatment

var	bonus0_mean	bonus0_sd	bonus1500_mean	bonus1500_sd	bonus500_mean	bonus500_sd
p0	0.553	0.265	0.573	0.248	0.530	0.251
job	0.760	0.430	0.805	0.399	0.829	0.379
myeduc	12.293	3.041	12.590	2.992	12.119	3.316
fyeduc	13.378	3.416	13.422	3.596	13.524	3.273
effort	19.549	9.460	18.303	10.592	18.477	10.475
math	5.476	1.468	5.388	1.258	5.386	1.360

2.2

Use the linear probability model to regress the dummy variable for completing all courses on the assignment of the three treatment groups. Interpret the treatment effects. Next include as additional regressors father's education, high-school math score and the subjective assessment about the pass probability.

For table 6, in the first model, we find that students who receive the 500 bonus are 0.007 percentage points more likely to pass the first year, and students who receive the f 1,500 bonus are 0.046 percentage points more likely to pass the first year, both relative to the group that receives no bonus. This effect is, however, not statistically significant, indicating that there is a high variance in passing within treatment groups, or, alternatively, that the sample size is too sample to statistically detect a relatively small effect size.

The point estimates change only slightly when including a vector of control variables, indicating that the treatment conditional on these (potential) confounders does not significantly increase the probability of passing. What we however find when we include the additional variables, is that the adjusted- R^2 increases

a lot, meaning a very high increase in explanatory power. These variables added also appear to be highly statistically significant.

2.3

Next also include as regressors in your model whether a student has a job and the amount of study effort. Comment on this approach. Do you consider this an improvement over (ii)?

The results are again displayed in table 6. Including job and study effort increases the adjusted- R^2 . The coefficient of P_0 becomes less insignificant, and effort is highly significant. This indicates that there is probably some correlation between P_0 and effort that was previously not captured by the model. Thus indicating an omitted variable bias problem beforehand. The rest of the coefficients and standard errors do not change that much by the additional variables. However, seeing that there is some indication of omitted variable bias, we conclude that this approach is an improvement.

```
attach(bonus)

# We need to regress pass on the treatment group assignment. This is
# a simple linear probability model. Model is denoted by prob_1.
# Note, we need to omit bonus0, cause of a dummy variable trap.
prob_1 <- lm(pass ~ bonus500 + bonus1500, data=bonus)
# Same as before, but now with some extra variables
prob_2 <- lm(pass ~ bonus500 + bonus1500 + fyeduc + math + p0, data=bonus)
# for q2.3:
prob_3 <- lm(pass ~ bonus500 + bonus1500 + fyeduc + math + p0 + job + effort, data=bonus)

# # q2.4; I think we should take the largest model. BAS???:
# prob_4 <- lm(pass ~ bonus500 + bonus1500 + fyeduc + math + p0 + job +
# effort + dropout + stp2001 + stp2004, data=bonus)

# Implement HC-robust standard errors (standard in linear prob model)
#library(sandwich)
#cov <- vcovHC(reg.model, type = "HC")
#robust.se <- sqrt(diag(cov))

# Create table:
stargazer(prob_1, prob_2, prob_3, header=FALSE, style='aer', label='tab:123')
```

2.4

Use your preferred model specification to estimate the effects of the financial incentives on some other outcomes: dropping out and credit points collected (in the first year and after three years).

% Table created by stargazer v.5.2.2 by Marek Hlavac, Harvard University. E-mail: hlavac at fas.harvard.edu
% Date and time: Fri, Jan 22, 2021 - 16:28:54

We have presented our results in table 7. We observe similar effects as before. It appears that math and effort remain highly significant

2.5

Given the sample size and the estimates you have obtained above, what would be the minimum detectable effect size of this experiment? In the next model, we observe that the point estimates for the

Table 6:

	pass		
	(1)	(2)	(3)
bonus500	0.007 (0.064)	0.015 (0.057)	0.023 (0.058)
bonus1500	0.046 (0.064)	0.048 (0.058)	0.056 (0.059)
fyeduc		-0.001 (0.007)	0.000 (0.007)
math		0.119*** (0.018)	0.124*** (0.018)
p0		0.249*** (0.095)	0.170* (0.098)
job			-0.062 (0.060)
effort			0.008*** (0.002)
Constant	0.195*** (0.045)	-0.576*** (0.128)	-0.678*** (0.144)
Observations	249	245	230
R ²	0.002	0.209	0.264
Adjusted R ²	-0.006	0.192	0.240
Residual Std. Error	0.411 (df = 246)	0.366 (df = 239)	0.358 (df = 222)
F Statistic	0.297 (df = 2; 246)	12.629*** (df = 5; 239)	11.359*** (df = 7; 222)

Notes:

*** Significant at the 1 percent level.

** Significant at the 5 percent level.

* Significant at the 10 percent level.

Table 7:

	dropout (1)	stp2001 (2)	stp2004 (3)
bonus500	−0.048 (0.068)	−0.144 (2.690)	−2.388 (7.999)
bonus1500	−0.057 (0.070)	0.227 (2.744)	0.294 (8.160)
fyeduc	0.009 (0.009)	−0.309 (0.337)	−1.231 (1.002)
math	−0.077*** (0.022)	6.565*** (0.856)	15.561*** (2.545)
p0	−0.150 (0.116)	13.582*** (4.567)	18.312 (13.579)
job	0.034 (0.071)	−0.904 (2.791)	5.628 (8.298)
effort	−0.018*** (0.003)	0.928*** (0.111)	2.920*** (0.331)
Constant	1.071*** (0.170)	−22.455*** (6.713)	−50.458** (19.962)
Observations	230	230	230
R ²	0.232	0.445	0.388
Adjusted R ²	0.207	0.427	0.369
Residual Std. Error (df = 222)	0.424	16.695	49.642
F Statistic (df = 7; 222)	9.563***	25.388***	20.112***

Notes:

*** Significant at the 1 percent level.

** Significant at the 5 percent level.

* Significant at the 10 percent level.

treatment effects increase a bit, indicating that effort was a confounder. Hence, we do consider this to be an improvement, although the treatment effects are still not significantly different from zero.

```
library(estimatr)
power <- lm_robust(pass ~ bonus500 + bonus1500 + fyeduc + math + p0 + job + effort, data=bonus)

# For now the constants:
n <- nrow(bonus)
df <- n - length(prob_3$coefficients)
alpha <- 0.05
t_95 <- qt(1-alpha/2, df)
sigma2 <- var(prob_3$residuals)

# Ik heb deze maar pre specified. We hebben twee verschillende treatments,
# zouden we dan niet twee power calculaties moeten doen?
p<-(83+84)/(82+83+84)

# How much power do we want? Range of powers:
q <- seq(from=0.6, to=0.9, by=0.1)

# Init empty vector of t values for the power:
t_q = rep(0, length(q))

# Fill vector iteratively and compute MDE for bonus500:
mde <- cbind(q*100, rep(0, length(q)))
counter <- 1
for (i in q){
  t_q[counter] <- qt(1-i, df) # A check: from slides, t_0.7 = -0.525...

  mde[counter,2] <- (t_95 - t_q[counter]) * sqrt( (1/(p*(1-p))) * (sigma2 / n) )

  counter <- counter + 1
}

kable(mde, col.names = c('Power (%)', 'MDE'))
```

Power (%)	MDE
60	0.1055720
70	0.1184614
80	0.1335612
90	0.1545456

2.6

Initially, the researchers were aiming at an increases in the pass rate of 10% points. How large should the sample size of the experiment have been in that case?

The proportion of treated subjects $p = \frac{83+84}{82+83+84} = 0.67$. Then in order to obtain the minimum size of the sample, we have re-written the general equation for the minimum detectable effect size in terms of n , the sample size. Formally, it takes the following form.

$$n = \frac{\sigma^2}{p(1-p)} / \left(\frac{mde}{(t_{1-\alpha} - t_{1-q})} \right)^2 \quad (1)$$

We have to however assume that we have a sufficiently large sample size that we can use the quantiles of the normal distribution. The reason is that the quantiles of the t-distribution depend on the number of degrees of freedom, which in turn depends on the sample size. Thus, we can not find an explicit expression in terms of n .

```
bonus_clean %>%
  group_by(kind_treatment) %>%
  summarize(test = n()) %>%
  kable()
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

kind_treatment	test
bonus0	82
bonus1500	83
bonus500	84

```
# Constants:
power_req <- 0.7
mde_req <- 0.1
n_95 <- qnorm(1-alpha/2)
n_q <- qnorm(1-power_req)

# Formula:
n_min <- (sigma2/(p*(1-p))) / (mde_req/(n_95-n_q))^2
```

Filling in the formula, we obtain that the sample size should be approximately 346.000.